

Back to Basics: Algebraic Data Types

I also do C++ training!
`arthur.j.odwyer@gmail.com`

Arthur O'Dwyer
2020-09-16

Outline

- Why the name “algebraic data types”? [3–18]
 - Memory layout diagrams. Why not `std::any`?
- Quick motivation for each type [19–27]
- Shared terminology [28–41] *Questions?*
- More about `optional` [42–49] *Questions?*
- More about `variant` [50–56] *Questions?*
- More about `pair` and `tuple` [57–69] *Questions?*

What do I mean by algebraic types?

`pair`

C++98. The original algebraic data type.

`tuple`

C++11.

`optional`

C++17.

`variant`

C++17, with minor tweaks to its constructors in C++20.

Why do we say “algebraic”?

It’s about the type’s number of possible values, a.k.a. the size of its *domain*. How many possible states might the object take on?

char	256 possible values
bool	2 possible values (true and false)
pair<char, bool>	$256 \times 2 = \mathbf{512}$ possible values
tuple<char, char, bool>	$256 \times 256 \times 2 = \mathbf{131072}$ possible values

Pair and tuple are product types

To find the size of the domain of a pair or tuple type, we take the *product* of the sizes of its element types.


A	A possible values
pair<A,B>	A×B possible values
tuple<A,B,C,...>	A×B×C×... possible values

Therefore pair and tuple are known as ***product types***.


Memory layout of a product type

The memory layout of pair or tuple is going to be pretty much the same as the layout of a plain old data struct.

`pair<A,B>`



`tuple<A,B,C>`



The compiler will do some **padding for alignment**, and may swap the order of the fields...


Memory layout of a product type

Here's the same two types as they would actually be laid out in memory by libc++...

`pair<A,B>`



`tuple<A,B,C>`




Memory layout of a product type

Here's the same two types as they would actually be laid out in memory by both libstdc++ and MSVC...

`pair<A,B>`




`tuple<A,B,C>`




Memory layout of a product type

Here's the same two types as they would be laid out in memory by a hypothetical library ABI. As far as I know, no vendor does this.

`pair<A,B>`




`tuple<A,B,C>`



pair appears as “basically” a struct

pair<A,B>



Notice how pair’s layout never changed.

The library specifies that pair is *basically* a simple struct:

```
template<class A, class B>
struct pair {
    A first;
    B second;
};
```

pair has two public data members, and first must be located at offset zero.

tuple has more freedom.

Why do we say “algebraic”?

Okay, that was product types. But there’s another mathematical operation we can use!

char	256 possible values
bool	2 possible values (true and false)
variant<char,bool>	$256 + 2 = \mathbf{258}$ possible values

tuple<char,bool> holds a char value **and** a bool value.

variant<char,bool> holds a char value **or** a bool value.

Variant is a sum type

To find the size of the domain of a variant type, we take the *sum* of the sizes of its alternative types.


A	A possible values
variant<A,B>	A+B possible values
variant<A,B,C,...>	A+B+C+... possible values

Therefore variant is known as a ***sum type***.


Memory layout of a sum type

Here's our *conceptual* picture of variant as opposed to tuple.

`tuple<A,B,C>`



`variant<A,B,C>`



Since we only need to store one alternative at a time, all three can be laid out at the *same* offset, just like in a union.

But this picture misses something!

Our ***conceptual*** picture isn't quite right. The STL's variant is designed to be type-safe.

A union doesn't know what type it holds. A variant does.


```
union U {  
    int i; float f;  
} u;  
u.f = 3.14f;  
int i = u.i;  
    // i gets 1078523331
```

```
std::variant<int, float> v;  
  
v = 3.14f;  
  
int i = std::get<int>(v);  
    // throws an exception!
```

Memory layout of a sum type

So, *actually*, variant also stores an “index” field.

`variant<A,B,C>`



The index field tells us which of the variant’s “alternatives” is currently active.

`size_t` which = `v.index()`;

`libc++`, `libstdc++`, and `MSVC` all use this same layout, with a 4-byte index field.

Finally, optional is another sum type

optional is the other *sum type*. Like pair, optional has been optimized for certain common use-cases.


An optional<A> can store any value of type A, or nullopt.

A	A possible values
variant<A,B>	A+B possible values
optional<A>	A+1 possible values


Memory layout of optional

`optional`, like `variant`, must store an “index” field. But its index field can be smaller — just a `bool` — and in practice that’s what we see.

`variant<A,C>`



`optional<C>`



There’s an implicit “or nullopt” alternative here; but the nullopt needn’t be stored physically in memory.

What about `std::any`?

- `std::any` also arrived in C++17, alongside `optional<T>` and `variant<Ts...>`.
- It is ***not*** an algebraic type. We can't use math to say anything interesting about its domain, which is “union of all copyable types.”
- `std::any` is a type-erasure type, similar to `std::function`.
- See my CppCon 2019 talk “[Back to Basics: Type Erasure](#).”

Still, I'll try to mention when there are commonalities between `std::any` and the algebraic types.

Quick motivation

Why use pair? (brief version)

pair is used many places in the classic STL.

Element type of map.

```
using M = std::map<int, int>;
```

Returned from insert.

```
M myMap = {{1,2}, {2,4}, {3,6}};
```

```
std::pair<int, int> item = *myMap.begin();
```

Returned from some algorithms,
like mismatch, equal_range,
and uninitialized_move_n.

```
std::pair<M::iterator, bool> ii = myMap.insert({2,3});
```

```
std::vector<int> a = {1,2,3,5,7};
```

```
std::vector<long> b = {1,2,3,4,5};
```

```
std::pair<decltype(a)::iterator, decltype(b)::iterator> mm =  
    std::mismatch(a.begin(), a.end(), b.begin(), b.end());
```

But std::ranges::mismatch returns
ranges::mismatch_result, not pair.

Why use tuple? (brief version)

Since the STL uses pair to simulate “returning multiple results,” you might imagine returning tuple when you need more than 2 results.

```
std::pair<int, int> minmax(int a, int b);  
std::tuple<int, int, int> minmidmax(int a, int b, int c);
```

You might also imagine using a tuple to simulate a “pack” of data members:

```
template<class R, class... Args>  
struct DeferredCall {  
    std::tuple<Args...> arguments_;  
    R operator()() const { ~~~ }  
};
```

However, C++20 moved toward returning dedicated class types, e.g.

```
struct minmax_result {  
    T min;  
    T max;  
};
```

and I recommend you do the same in your code.

Why use tuple? (brief version)

tuple is used (arcanelly) to forward sets of arguments to pair's constructor.
More on this later.

```
std::pair<std::string, std::string> myPair(  
    std::piecewise_construct,  
    std::make_tuple("abc", 3),  
    std::make_tuple(100, '*')  
);
```

Construct `myPair.first`
as `string("abc", 3)`
and `myPair.second` as
`string(100, '*')`.

We'll also see a couple of useful idioms with `std::tie`.

Why use optional? (brief version)

`optional` is never used in the STL, but it's the most useful for your own code.

In a return type, it represents “no answer.”

```
std::optional<std::string> oGetenv(const char *name) {  
    if (name is present) {  
        return std::string(variable's value);  
    } else {  
        return std::nullopt;  
    }  
}
```

Why use optional? (brief version)

In a variable, it represents “no setting” (if that’s distinct from “set to empty”).

```
bool hasIncludePath_ = false;
std::string includePath_;
std::string getIncludePath() const {
    return hasIncludePath_ ? includePath_ : "/usr/include";
}
```

```
std::optional<std::string> includePath_;
std::string getIncludePath() const {
    return includePath_.value_or("/usr/include");
}
```


Why use optional? (brief version)

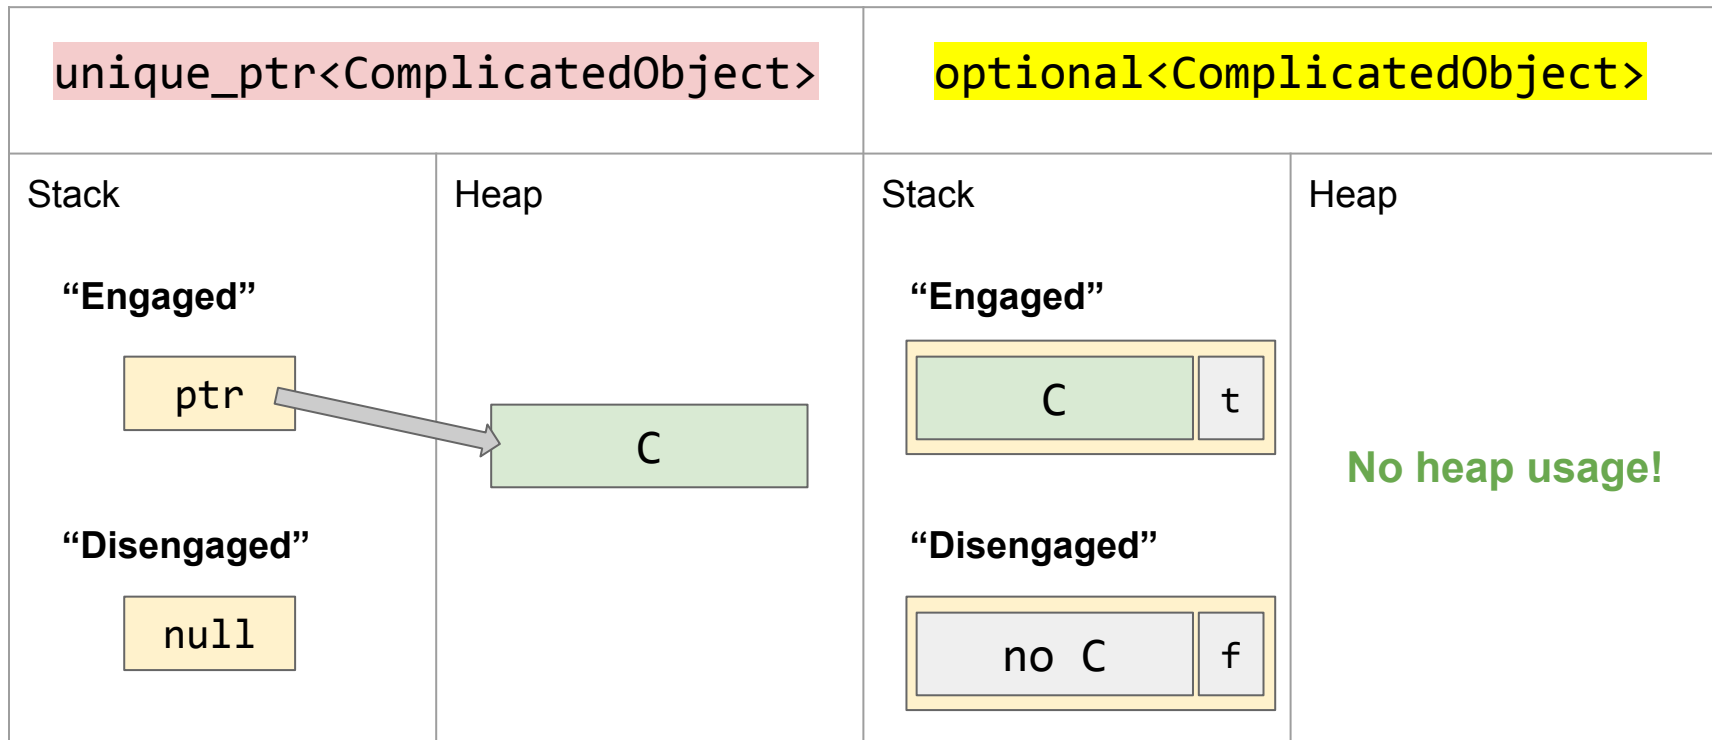
Or, it represents “not initialized yet,” for types with no default constructor. This gives us the benefits of *dynamic lifetime*, without any *heap usage*.

```
std::unique_ptr<ComplicatedObject> obj_ = nullptr;
void setComplicated(int a, int b) {
    obj_ = std::make_unique<ComplicatedObject>(a, b);
}
```

```
std::optional<ComplicatedObject> obj_ = std::nullopt;
void setComplicated(int a, int b) {
    obj_.emplace(a, b);
}
```

The next slide draws this out in a diagram.

Why use optional? (brief version)



Why use variant? (brief version)

When you need a “sum type,” you might directly use variant.

Or you might use it as an implementation detail in writing your own sum type.

```
struct Simple { int fromx, fromy, tox, toy; };  
struct Capture { int fromx, fromy, tox, toy; bool enpassant; };  
struct Castle { bool kingside; };  
  
class ChessMove {  
    std::variant<Simple, Capture, Castle> details_  
  
    bool isSimple() const { return details_.index() == 0; }  
    bool isCapturing() const { return details_.index() == 1; }  
    bool isCastling() const { return details_.index() == 2; }  
};
```

Common features of algebraic types

Now that you know how they look in memory, and have some general motivation for each of them, let's talk about the C++ side of things.

All four algebraic types — `pair`, `tuple`, `optional`, `variant` — use some common concepts and vocabulary.

Let's start with the new vocabulary.


“Engaged” and “disengaged”

- An `optional` which holds a value is often said to be *engaged*.
- An empty `optional` is said to be *disengaged*.
- These terms aren’t relevant to `pair`, `tuple`, or `variant`, but I explain them because they’re useful.
- Test whether an `optional` is engaged with `o.has_value()`, or via (explicit or contextual) conversion to `bool`.

```
if (auto o = oGetenv("foo")) ...  
if (o) ...
```

“Engaged” and “disengaged”

By the way... Types which own a unique resource without being semantically “pointer-like” (and thus “nullable”) tend to have some “disengaged” state. This includes type-erased owners:

- `std::any` — `a.has_value()`
 - `std::function` — conversion to `bool`
- 
- optional supports both of these methods!*

And other movable RAI types:

- `std::future` — `fut.valid()`
- `std::unique_lock` — `(lk.mutex() != nullptr)`
- `std::promise`, `std::thread` — no explicit way to test

“Emplace” and in-place construction

- A variant or optional can be thought of as a “buffer” that *may or may not* hold a value of some given type.
- We can explicitly **emplace** an object into that buffer, in the same way as we might `emplace` a new element into a vector or set.
- The old value is destroyed before the new value is constructed.
- The arguments are perfectly forwarded to T’s constructor.

```
vectorOfString.emplace_back("abc", 3);
```

```
optionalString.emplace("abc", 3);
```

```
variantOfStringAndInt.emplace<std::string>("abc", 3);
```

“Emplace” and in-place construction

To construct a variant or optional with an object already emplaced, you have two options:

- The “basic” option: `variant` and `optional` are implicitly convertible from their alternative type(s).

```
optional<string> o1 = "foo";
```

```
variant<int, string> v1 = "foo";
```

// Generally speaking, these will work fine.

// But maybe not if the stored type is non-movable!

“Emplace” and in-place construction

- The “advanced” option: `variant` and `optional` have “in-place constructor” overloads to construct their contents directly in-place.

```
optional<string> o2(std::in_place, "abc", 3);
```

```
variant<int, string> v2(  
    std::in_place_type<std::string>, "abc", 3);
```

```
variant<int, string> v3(  
    std::in_place_index<1>, "abc", 3);
```

```
Also: std::any a(std::in_place_type<std::string>, "abc", 3);  
      a.emplace<std::string>("abc", 3);
```

std::get maps an index to a value

- You can think of a pair, tuple, or array as a **mapping** from an index in the range [0..n) to a value.
- Use `std::get<index>(obj)` to extract the *index*'th value.

```
std::pair<int, bool> p = {42, true};  
std::cout << std::get<0>(p); // synonym for p.first  
std::get<1>(p) = false;      // synonym for p.second
```

```
std::array<int, 10> arr;  
std::get<4>(arr) = 42;        // synonym for arr[4]
```

```
std::tuple<int, bool, double> t;  
std::cout << std::get<1>(t); // the bool element of t
```

`std::get` maps an index to a value

- `variant` is similar, but it's only a partial mapping.
- `std::get<I>` for any `I` other than the active index will throw an exception of type `bad_variant_access`. `variant` remains type-safe.

```
variant<int, string> v = 42;
```

```
std::get<0>(v) = 43;      // OK, it's the active alternative
```

```
std::get<1>(v) = "def";   // Throws std::bad_variant_access
```

```
std::get<2>(v) = 1;       // Error: get<2>(v) is just ill-formed
```

std::get can also map *type* to value

For pair, tuple, and variant, std::get also works with a type parameter — as long as it's unique and unambiguous.

```
pair<int, string> p = {42, "abc"};  
std::cout << std::get<int>(p);    // synonym for p.first  
std::get<string>(p) = "def";      // synonym for p.second
```

The implementation, simply, is that get has two overloads:

```
template<size_t I, class... Ts> T& get(tuple<Ts...>&);  
template<class T, class... Ts> T& get(tuple<Ts...>&);
```

Easy to #include and declare

Or, using CTAD and/or direct-initialization,
but personally I don't recommend either...

```
#include <utility>
```

```
std::pair<int, char> p = {9, 'W'};
```

```
std::pair p{9, 'W'};
```

```
#include <tuple>
```

```
std::tuple<int, char> t = {9, 'W'};
```

```
std::tuple t{9, 'W'};
```

```
#include <optional>
```

```
std::optional<int> o = 9;
```

```
std::optional o{9};
```

```
std::optional<int> no = std::nullopt;
```

```
std::optional<int> o{};
```

```
#include <variant>
```

```
std::variant<int, char> v = 'W';
```

```
variant<int, char> v{'W'};
```

Recursive default-construction

```
std::pair<std::string, int> p1{}, p2;
```

Both construct the pair with `{ "", 0 }`, as if by

```
std::string first{};  
int second{};
```

```
std::tuple<int, char> t1{}, t2;
```

Likewise, constructs the tuple with `{ 0, 0 }` in both cases.

```
std::optional<int> o1{}, o2;
```

In both cases, equivalent to `nullopt`: optionals are default-constructed into their disengaged state.

```
std::variant<int, char> v1{}, v2;
```

A default-constructed variant value-initializes its 0'th alternative. In this case the `int` gets `0`.

Special member functions

`pair`, `tuple`, `variant`, and `optional` all inherit their special members from their constituent types in the “natural” way.

- `foo<Ts...>` is copy-constructible if all `Ts...` are copy-constructible.
- `foo<Ts...>` is copy-assignable if all `Ts...` are copy-assignable.
 - `variant` also requires that all `Ts...` be copy-constructible, since it might need to change the active member of the LHS.
- `foo<Ts...>` is default-constructible if all `Ts...` are default-constructible.
 - Exception: Default-constructing a `variant<A,B,C>` will default-construct only its `A` alternative; so only `A` needs to be default-constructible.
 - And `optional` defaults to disengaged, so it's always default-constructible.

Comparison operators

Finally, `pair`, `tuple`, `variant`, and `optional` all have “natural” comparison operators.

- `pair<Ts...>` and `tuple<Ts...>` are ordered sequences, so they’re compared lexicographically.
- `optional<T>` is compared just like `T` would be, except that `nullopt` compares less-than any value from `T`’s domain.
- `variant<Ts...>` can be thought of as an ordered sequence `{index(), value}`, compared lexicographically.
 - Or, equivalently, `variant<Ts...>` can be thought of as `tuple<optional<Ts>...>` and compared lexicographically on *that* basis.

High-level questions?

FYI, we're about to discuss idioms around optional, variant, and pair/tuple, in that order.

About optional

```
#include <optional>
```

```
std::optional<int> o = std::nullopt;
```

optional for “not (yet) set”

Recall that `optional<T>` means “either a T, or nothing.”

It’s commonly used in business logic:

```
struct NetworkConnection {  
    std::optional<std::string> password_  
    std::optional<Certificate> cert_  
};
```

`optional<string>` gives us distinct states for “no password” and “password is the empty string.”

`optional<Certificate>` gives us a state for “no certificate provided.”

Using optional

```
Certificate NetworkConnection::getCert() const {  
    if (cert_.has_value()) {  
        return cert_.value();  
    } else {  
        return getDefaultCertificate();  
    }  
}
```

`optional::value()`
retrieves the held value
(propagating value
category appropriately).

If disengaged,
`value()` will throw
`bad_optional_access`.

Technically this means it
does a redundant check in
this case. But inlining will
save us.

To eliminate that check...

Using optional

```
Certificate NetworkConnection::getCert() const {  
    if (cert_) {  
        return *cert_;  
    } else {  
        return getDefaultCertificate();  
    }  
}
```

For clarity, I would always write `.has_value()`!

The one thing `operator bool` lets you do is write code that works generically with both `T*` and `optional<T>`, as part of a gradual upgrade strategy.

`optional` has an explicit conversion to `bool`, synonymous with `has_value()`.

`operator*` is like `.value()`, except that it has UB instead of throwing, so it can skip the check.

Mnemonic: Just like `vector`'s `operator[]` and `.at()`, the terse punctuation has UB; the named method throws.

Using `optional::value_or`

```
Certificate NetworkConnection::getCert() const {  
    return cert_.value_or(  
        getDefaultCertificate()  
    );  
}
```

You can hide the test inside this library-provided convenience method.

Watch out for side effects, though. In the previous version, we never called `getDefaultCertificate` unless the optional was disengaged. In this version, we call `getDefaultCertificate` *before* testing the optional.

Setters for optional fields

```
class NetworkConnection {  
    std::optional<Cert> cert_  
  
    void setCert(???);  
};
```

For setters, do this!



- `void setCert(std::optional<Cert>)`
 - Usually what you want.
- `void setCert(Cert&&) + void setCert(std::optional<Cert>)`
 - Might save you a move, when called as `o.setCert(Cert("foo"))`.
- `void setCert(const std::optional<Cert>&)`
 - Basically never what you want. The caller usually can't give you a reference to "their" optional object because they haven't got one, so you'll get a temporary.

optional for “optional parameters”?

```
void openConnection(std::string_view host,  
                   std::optional<Cert> cert = std::nullopt);  
  
    openConnection("example.com", Cert("foo")); // OK, works  
    openConnection("example.com");              // OK, works  
    openConnection("example.com", std::nullopt); // OK, works
```

This may be exactly what you’re looking for. Personally, I wouldn’t use default function arguments for this (or anything ever). I’d either expect the caller to be passing along an `optional<Cert>` they got from somewhere else—

```
void openConnection(std::string_view, const std::optional<Cert>&);
```

or I’d expect the two different signatures to be used in different situations—

```
void openConnectionWithCert(std::string_view, const Cert&);  
void openConnectionWithoutCert(std::string_view);
```


Questions on optional?

About variant

```
#include <variant>
```

```
std::variant<int, double> v = 3.14;
```

Detect the active alternative

There are several ways to figure out which alternative of a variant is active.

- `v.index()` — returns the active index as a `size_t`
- `std::holds_alternative<int>(v)` — I'd never recommend this
- `std::get_if<0>(&v)` — returns a pointer to the specified alternative, if it's active; otherwise returns `nullptr`
- `std::get_if<int>(&v)` — same but with types

Visitation

```
std::variant<int, double, std::string> v;  
  
if (v.index() == 0) {  
    std::cout << std::get<int>(v) << "\n";  
} else if (v.index() == 1) {  
    std::cout << std::get<double>(v) << "\n";  
} else if (v.index() == 2) {  
    std::cout << std::get<std::string>(v) << "\n";  
}
```

This code works,
but it's ugly and
error-prone.

We could have
used `std::get<0>`
etc. I'm
foreshadowing the
next slide.

Visitation

```
std::variant<int, double, std::string> v;

if (int *pi = std::get_if<int>(&v)) {
    std::cout << *pi << "\n";
} else if (double *pd = std::get_if<double>(&v)) {
    std::cout << *pd << "\n";
} else if (auto *ps = std::get_if<std::string>(&v)) {
    std::cout << *ps << "\n";
}
```

`std::get_if` is like `std::get`, but

- it takes a pointer
- it returns a pointer
- it returns `nullptr` if the specified alternative isn't active.

Compare it with `std::any_cast`.

Visitation with `std::visit`

```
std::variant<int, double, std::string> v;  
  
auto printme = [](const auto& x) {  
    std::cout << x << "\n";  
};  
  
std::visit(printme, v);
```

When every branch of the visitation does “the same thing” (syntactically speaking), you can use `std::visit`.

Internally, it does ***exactly the same thing*** as on the previous slide (modulo some clever optimizations). It branches on `v.index()`, and it instantiates a call to `printme(a)` for each alternative `a` that might be held by the variant.

Standard library functions usually take the lambda as their *last* argument, but `visit` takes it *first*, to leave room for a variadic number of variants.

With a two-parameter lambda `printus`, you could call `std::visit(printus, v1, v2)`.

valueless_by_exception

```
std::variant<std::unique_ptr<int>, std::string> v;  
  
v = std::make_unique<int>(127);  
std::string x = "long enough to require heap allocation";  
try {  
    v.emplace<1>(x);  
} catch (const std::bad_alloc&) {  
    assert(v.valueless_by_exception());  
    assert(v.index() == size_t(-1));  
}
```

`v.emplace<1>(x)` destroys the `unique_ptr` to make room for the string. Contrariwise, `v=x` will effectively “copy-and-swap,” leaving `v` with its old value, as long as `x` is copyable and/or nothrow-movable.

This state basically never happens unless you use `emplace` in an unwise manner; and even if it does, you should keep it confined to catch blocks.

Questions on variant?

About pair and tuple

```
#include <utility>      // for pair
#include <tuple>         // for tuple
#include <functional>    // for reference_wrapper

std::pair<int, int> p = {1,2};
std::tuple<int, double, int> t = {1,2,3};
```

Making pairs and tuples

Pairs and tuples represent “sequences” and so it makes sense to initialize them with braced initializer lists where convenient:

```
std::pair<int, bool> foo() {  
    std::tuple<int, int, int> t = {1,2,3};  
    return {42, false};  
}
```

C++17 CTAD alert!

Or, use these helper functions:

```
auto foo() {  
    std::tuple t = std::make_tuple(1, 2, 3);  
    return std::make_pair(42, false);  
}
```

Function templates `make_pair` and `make_tuple` will deduce their argument types.

Distinguish from `make_optional<T>`, `make_any<T>`, `make_shared<T>`, `make_unique<T>`.

pairs and tuples of references

You can't make an `optional<T&>` or `variant<T&, U&>`, but you **can** make a `tuple<T&, U&, ...>`, which will have “assign-through” behavior:

```
int a=1, b=2, c=3, d=4;
std::tuple<int&, int&> ab {a, b};
std::tuple<int&, int&> cd {c, d};

ab = cd; // assign cd's values “through” ab

assert(a == 3 && b == 4);
```

Making tuples of references

You can implicitly create a tuple of references in at least three ways:

```
int a=1, b=2;
```

```
auto ab1 = std::make_tuple(std::ref(a), std::ref(b));
```

```
auto ab2 = std::tie(a, b);
```

```
auto ab3 = std::forward_as_tuple(a, b);
```

```
static_assert(std::is_same_v<decltype(ab1), std::tuple<int&, int&>>);
```

- `make_tuple` always captures values, except that it “decays” `reference_wrapper<T>` into `T&`.
- `tie` purposely works only on lvalues, and captures lvalue references.
- `forward_as_tuple` always captures references (to either lvalues or rvalues).

forward_as_tuple for argument lists

Remember our discussion of in-place construction?

```
std::optional<std::string> o(std::in_place, otherString, len);
```

For pair and tuple, we have to give **two** constructor argument lists.

So instead of an “in-place” constructor, we have a “piecewise” constructor.

```
std::pair<std::string, std::string> p(  
    std::piecewise_construct,  
    std::forward_as_tuple(otherString, len),  
    std::forward_as_tuple(std::move(thirdString))  
);
```

This is the primary use-case for forward_as_tuple. Very niche.

Multiple assignment with tie

We can use “assign-through” to simulate multiple assignment.

```
using Set = std::set<int>;
```

```
Set numbers;
```

```
Set::iterator it;
```

```
bool inserted;
```

```
std::tie(it, inserted) = numbers.insert(n);
```

tie returns tuple<iterator&, bool&>



insert returns pair<iterator, bool>




The assignment operator for tuple<iterator&, bool&> “assigns-through,” updating our named variables on the LHS with the values of the pair on the RHS.

Multiple assignment with tie

The STL provides a special tag type for “ignoring” one field during assign-through.

```
std::tuple<int, int, int> getXYZ();  
int x, y;  
std::tie(x, y, std::ignore) = getXYZ();
```

tie returns tuple<int&, int&, const Magic&>



std::ignore is a global variable of some magic type whose operator= accepts any type on the RHS and swallows it with no effect.

Multiple assignment with tie

Warning! Assign-through is not the same as *simultaneous* multiple assignment, as found in languages such as Python and Perl.

```
int x, y;
```

```
// Exchange the values of x and y
```

```
std::tie(x, y) = std::make_tuple(y, x);
```

```
// Exchange the values of x and y? No!
```

```
std::tie(x, y) = std::tie(y, x);
```

```
std::tie(x, y) = {y, x};
```

Don't try to be "clever" by over-using `std::tie`, and you'll never care about this particular pitfall.

Write one assignment per line. In this specific example, we should have used `std::swap`!

Using tie for comparison

The other useful idiom is “compare-through.”

Remember, tuple (and pair) do lexicographical comparison.

```
class Name {  
    std::string firstname, lastname;  
  
    std::tuple<const std::string&, const std::string&> asKey() const {  
        return std::tie(lastname, firstname);  
    }  
  
    friend bool operator<(const Name& a, const Name& b) {  
        return a.asKey() < b.asKey();  
    }  
};
```

I wrote out this return type for clarity.
In real life you'd just say auto.

tuple is not great for public APIs

Let's say I have a function that returns a hostname, a cert, and a time-to-live.

```
auto generateDefaultCert()  
    -> std::tuple<std::string, Cert, double>;
```

// C++ before '17

```
auto hct = generateDefaultCert();  
std::cout << "Made cert for host " << std::get<0>(hct) << "\n";
```

// C++17 structured binding

```
auto [host, cert, ttl] = generateDefaultCert();  
std::cout << "Made cert for host " << host << "\n";
```

tuple is not great for public APIs

Using a named class type, with named fields, is *much* friendlier.

```
struct CertInfo { std::string host; Cert cert; double ttl; };  
CertInfo generateDefaultCert();
```

// C++ before '17

```
auto info = generateDefaultCert();  
std::cout << "Made cert for host " << info.host << "\n";
```

// C++17 structured binding still permits this

```
auto [host, cert, ttl] = generateDefaultCert();  
std::cout << "Made cert for host " << host << "\n";
```

In conclusion

- Use `std::optional` for “maybe a T” or “not a T **yet**”
- Use `std::pair` and `std::tuple` for implementation details
 - Such as `std::tie` for “multiple assignment”
 - But prefer named classes in public interfaces
- Remember `std::in_place` and `std::piecewise_construct`
- Remember `std::visit` for variants
- Forget `std::variant::valueless_by_exception`

Questions?

Bonus Slides

Why use variant? (brief version)

variant can also be used as a poor man's Expected<T>.

I don't recommend this, but it shows basically what variant is capable of.

```
std::variant<std::string, std::errc> vGetenv(const char *name);

if (auto v = vGetenv("foo"); std::get_if<std::string>(&v)) {
    const auto& value = std::get<std::string>(v);
    std::cout << "Value is: " << value << "\n";
} else {
    std::error_condition error = std::get<std::errc>(v);
    std::cout << "Error was: " << error.message() << "\n";
}
```

I associate value_or with this pitfall

```
class NetworkConnection {  
    static const int defaultIdleTimeout = 1000;  
    ~~~  
};
```

Can you spot the bug?

Solution on next slide.

```
int NetworkConnection::getTimeout() const {  
    return idleTimeout_.value_or(defaultIdleTimeout);  
}
```


```
test.o: In function `NetworkConnection::getTimeout() const':  
test.cpp:10: undefined reference to `NetworkConnection::defaultIdleTimeout'
```


I associate value_or with this pitfall

```
class NetworkConnection {  
    static constexpr int defaultIdleTimeout = 1000;  
    ~~~  
};  
  
int NetworkConnection::getTimeout() const {  
    return idleTimeout_.value_or(defaultIdleTimeout);  
}
```

Static const members
must still have an
out-of-line definition!

Static constexpr
members, or C++17
static *inline* const
members, are freed of
that restriction.

value_or takes its parameter by forwarding reference, which means it wants defaultIdleTimeout's address. Often value_or is the only place in the program that asks for the *address*, rather than the *value*, of a static const data member.